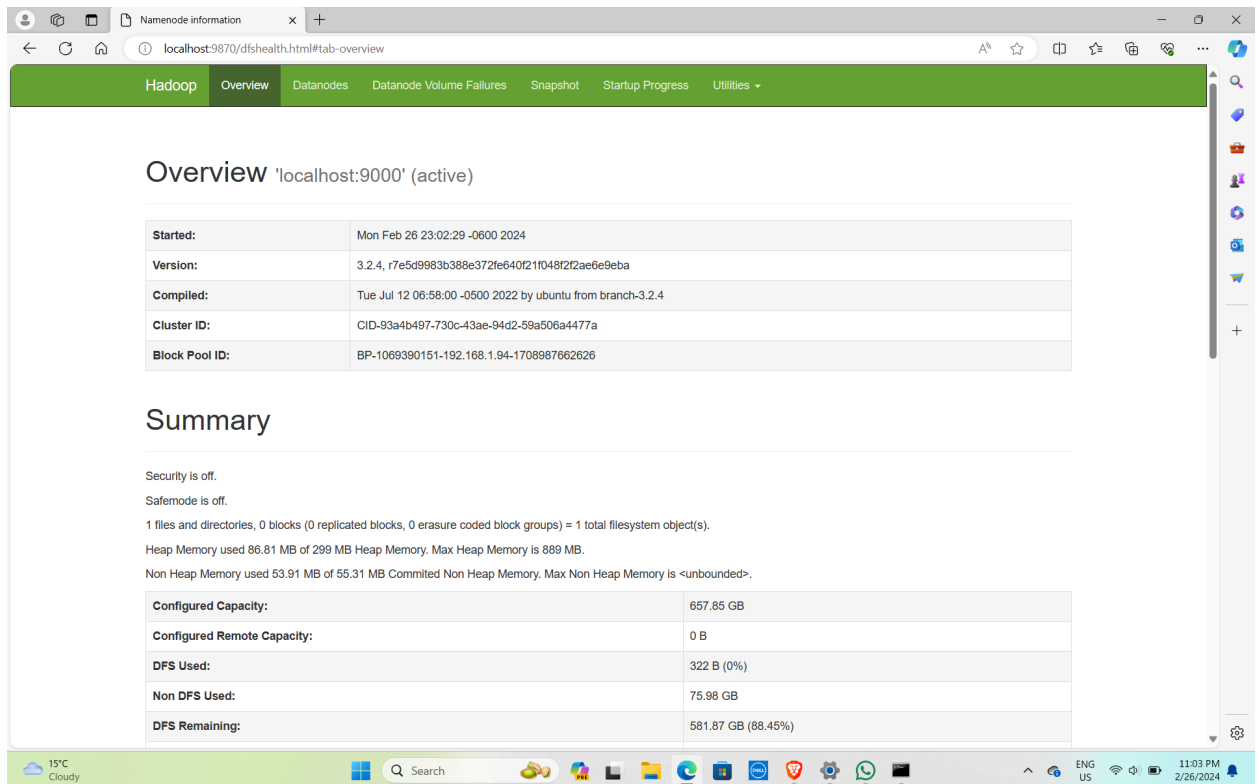


Day21 - Feb26th 2024

1. Woke up at 6 and completed all my things
2. Cooked food and headed to library with my friends
3. Started marketing my profile for data engineering ..but all I receive are rejection mails
4. Opened leetcode and solved one medium binary search question ..was not able to find the optimal logic..so referred a solution
5. Took me around 2 hrs to debug the code and to analyze the time complexity..but soon I'll get improved

6. Completed installing java and hadoop on my local disk to start pyspark project



The screenshot shows the Hadoop NameNode Overview page for 'localhost:9000' (active). The page has a green header with tabs: Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The Overview tab is selected.

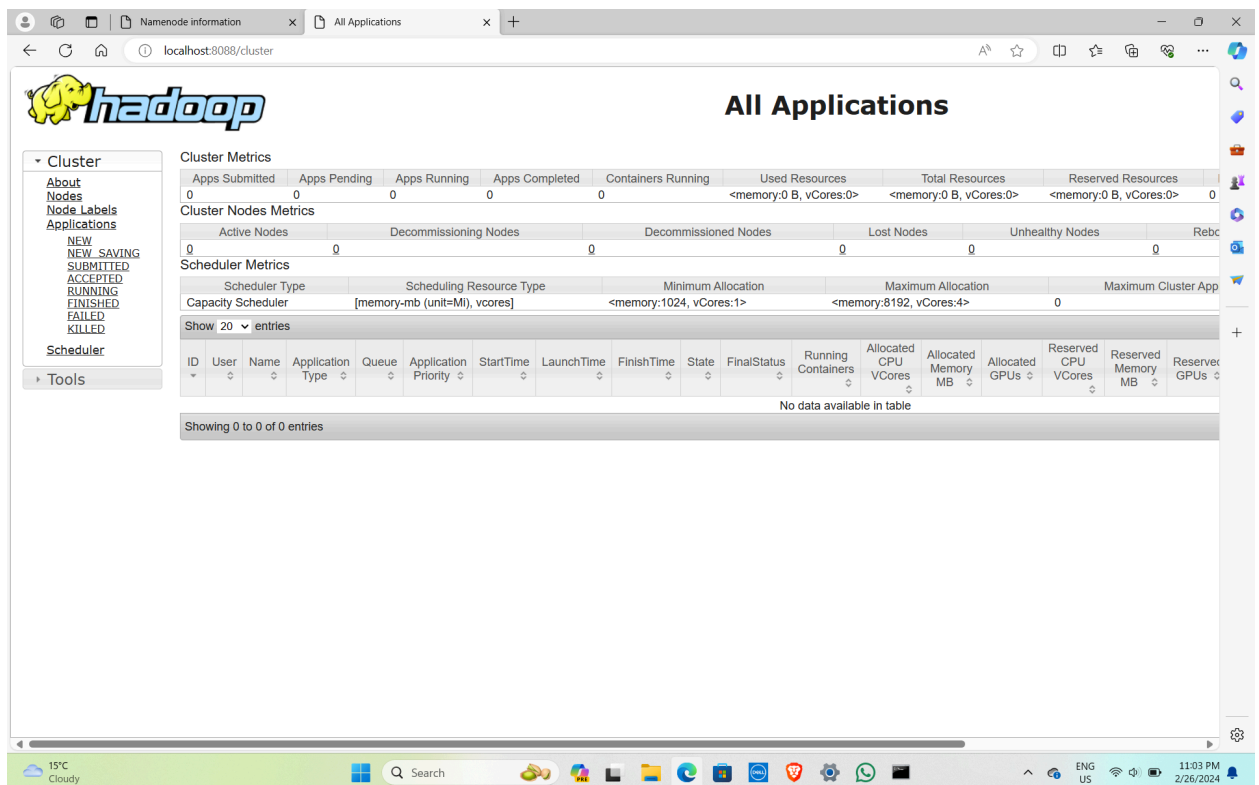
Overview 'localhost:9000' (active)

Started:	Mon Feb 26 23:02:29 -0600 2024
Version:	3.2.4, r7e5d9983b388e372fe640f21f048f2f2ae6e9eba
Compiled:	Tue Jul 12 06:58:00 -0500 2022 by ubuntu from branch-3.2.4
Cluster ID:	CID-93a4b497-730c-43ae-94d2-59a506a4477a
Block Pool ID:	BP-1069390151-192.168.1.94-1708987662626

Summary

Security is off.
Safemode is off.
1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 86.81 MB of 299 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 53.91 MB of 55.31 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	657.85 GB
Configured Remote Capacity:	0 B
DFS Used:	322 B (0%)
Non DFS Used:	75.98 GB
DFS Remaining:	581.87 GB (88.45%)



The screenshot shows the Hadoop NameNode All Applications page for 'localhost:8088/cluster'. The page has a green header with tabs: Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The Overview tab is selected.

All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Used Resources	Total Resources	Reserved Resources
0	0	0	0	0	<memory:0 B, vCores:0>	<memory:0 B, vCores:0>	<memory:0 B, vCores:0>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Reb...
0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster App
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs	Reserved CPU V-Cores	Reserved Memory MB	Reserved GPUs
No data available in table																	

Showing 0 to 0 of 0 entries

7. Ended my day by solving two complex SQL question from Ankit's YT

SQLQuery1.sql - KAUSHI\SQLEXPRESS.master (KAUSHI\jamka (64)) - Microsoft SQL Server Management Studio

Object Explorer

Connect - KAUSHI\SQLEXPRESS (SQL Server 16.0.100.100)

SQLQuery1.sql - K...KAUSHI\jamka (64)*

```
select * from company_users;

--Q: Find companies who have atleast 2 users who speaks

with cte as(
select company_id,user_id,count(language) as 1 from company_users
where language in('English','German')
group by company_id,user_id
having count(language) = 2)
select company_id
from cte
group by company_id
having count(user_id) >= 2

/* Explanation:

Step 1: First we filter the table with people speaking only English and German
Step 2: Now we use group by on company and user to find count of people
       Speaking two languages and framed it in CTE
Step 3: Next we will find company where count of users >=2
```

Results

	company_id	user_id	language
1	1	1	English
2	1	1	German
3	1	2	English
4	1	3	German
5	1	3	English
6	1	4	English
7	2	5	English
8	2	5	German
9	2	5	Spanish
10	2	6	German
11	2	6	Spanish
12	2	7	English

Query executed successfully.

KAUSHI\SQLEXPRESS (16.0 RTM) KAUSHI\jamka (64) master | 00:00:00 12 rows

Ready 15°C Mostly cloudy

8.

SQLQuery1.sql - KAUSHI\SQLEXPRESS.master (KAUSHI\jamka (64)) - Microsoft SQL Server Management Studio

Object Explorer

Connect - KAUSHI\SQLEXPRESS (SQL Server 16.0.100.100)

SQLQuery1.sql - K...KAUSHI\jamka (64)*

```
select * from company_users;

--Q: Find companies who have atleast 2 users who speaks

with cte as(
select company_id,user_id,count(language) as 1 from company_users
where language in('English','German')
group by company_id,user_id
having count(language) = 2)
select company_id
from cte
group by company_id
having count(user_id) >= 2

/* Explanation:

Step 1: First we filter the table with people speaking only English and German
Step 2: Now we use group by on company and user to find count of people
       Speaking two languages and framed it in CTE
Step 3: Next we will find company where count of users >=2
```

Results

	company_id
1	1

Query executed successfully.

KAUSHI\SQLEXPRESS (16.0 RTM) KAUSHI\jamka (64) master | 00:00:00 1 rows

Ready 15°C Mostly cloudy